

Use Case Story: Building a RAG (Retrieval-Augmented Generation) Chatbot

Objective:

Build a **Retrieval-Augmented Generation (RAG) Chatbot** that allows users to interact with documents (PDF, audio, video) via a simple web UI and get responses with proper source references.

Use Case Overview:

1. The user accesses a **Streamlit-based UI**.
2. The user uploads a **PDF, Audio, or Video file**.
3. The file is processed and indexed into a vector store.
4. The user can then chat with the bot by asking questions related to the uploaded files.
5. The chatbot generates a relevant response using an **open-source LLM model** and shows the source reference.

Functional Requirements:

1 *User Interface (Streamlit):*

- Simple and clean UI where the user can:
 - Upload a file (PDF, Audio, Video).
 - Enter a question in a chat window.
 - See the chatbot's response along with source reference.

2 *Backend API (FastAPI):*

- Expose endpoints for:
 - File upload (PDF, Audio, Video).
 - Query handling.

- Implement middleware to:
 - Log incoming requests and responses.
 - Handle errors gracefully.
- Maintain basic logs (request timestamps, input size, user query).

LLM Integration:

- Use any open-source model (e.g., Hugging Face models).
- When a user asks a question:
 - Retrieve relevant context from the uploaded documents (using vector search).
 - Generate an answer based on the retrieved context.
 - Attach source reference in the answer.

Docker Setup:

- Provide a **Dockerfile** and **docker-compose.yml**.
- Ensure the entire app (Streamlit UI + FastAPI + LLM) runs in containers.
- Easy startup with a single command like:

`docker-compose up -build`

Project Management with UV:

- Use **UV (Universal Virtual environment)** for:
 - Dependency management (Python packages).
 - Environment isolation.
 - Project structure management.
 - Python version - 3.11

Example User Flow:

1. Open the app in the browser.
2. Upload a PDF document (e.g., product manual).

3. Ask: “What are the safety precautions for using this device?”
4. Chatbot responds:

“the device must not be used near water due to electrical risks.

Reference : Page 1 – pdf name”

Non-Functional Requirements:

- Easy-to-use UI.
- Lightweight and modular backend structure.
- Fast response time (within seconds).
- Reproducible environment via Docker.

Goal for Interns:

- Design the Streamlit frontend to upload files and interact with the chatbot.
- Build a robust FastAPI backend with middleware and logging.
- Integrate an open-source LLM model for generating responses.
- Implement vector-based retrieval from uploaded files.
- Containerize the entire app using Docker.